Hamilton Filtering for Market Regime Identification

Detecting Structural Breaks and Spurious Persistence in S&P 500

Alfredo Berlose
MSc in Quantitative Finance and Risk Management
Ca' Foscari University of Venice

April 2026

Abstract

This paper investigates the dynamic behavior of S&P 500 index returns (1998–2026) by employing a Markov-Switching regression framework. Traditional econometric models often assume constant parameters, leading to "spurious persistence" where structural breaks are misinterpreted as long-memory processes. Drawing on the theory of non-constant parameter models, we demonstrate that financial volatility is better characterized by discrete regime shifts rather than deterministic linear dependence. By implementing an adaptive Hamilton Filter, we derive real-time filtered probabilities to map systemic fragility. The empirical results, estimated via Quasi-Maximum Likelihood (QML), reveal a sharp dichotomy between a stable "Bull" regime and a high-variance "Crisis" regime, where the latter exhibits a variance approximately 7.7 times higher than the former. Furthermore, out-of-sample validation confirms that the regime-switching approach provides superior risk diagnostics compared to fixed-parameter ARMA benchmarks, particularly during tail events and periods of heightened market stress.

1 Introduction

The identification of structural shifts and speculative dynamics represents a cornerstone of modern financial econometrics. Standard linear models, while computationally efficient, often rely on the assumption of parameter stability over long horizons. However, as noted in the pedagogical framework of "Models with Non-Constant Parameters," financial time series are frequently characterized by discrete breaks in mean and variance, which can lead to misleading statistical inferences.

In particular, the phenomenon of "spurious persistence" arises when a model fails to account for regime switches. An econometrician might observe high autocorrelation in squared returns—suggesting a long-memory ARCH/GARCH process—when, in reality, the series has undergone a localized shift in its volatility level. This paper addresses this challenge by implementing a Markov-Switching (MS) framework on the S&P 500 index, moving beyond static diagnostics toward a dynamic mapping of market fragility.

2 Literature Review: From Explosive Roots to Adaptive Filtering

The detection of speculative bubbles has traditionally relied on unit root testing. The Augmented Dickey-Fuller (ADF) test serves as the baseline for stationarity; however, it lacks the power to identify "periodically collapsing bubbles." As discussed by Phillips et al. (2015), advanced recursive tests like the Generalized Supremum ADF (GSADF) were developed to date-stamp multiple explosive episodes *ex-post*.

While GSADF provides a rigorous historical account of bubbles, it presents two major limitations for real-time risk management:

1. **Computational Burden:** The recursive nature of GSADF requires repeated re-estimations as new data points arrive.

2. **Binary Output:** These tests typically provide a "yes/no" signal regarding explosiveness rather than a probabilistic measure of state transition.

Following the approach of Habibi (2026), this study utilizes the Hamilton Filter as a superior alternative for real-time monitoring. By treating the market state as a latent variable S_t, we derive time-varying conditional probabilities of being in a "Crisis" regime. This methodology aligns with the principles of "Financial Econometrics: Realized Volatility and Inference," shifting the focus from retrospective identification to adaptive, forward-looking diagnostic mapping.

3 Econometric Framework and Mathematical Derivation

The cornerstone of our analysis is the *Markov-Switching Regression* model, which treats market dynamics as a stochastic process governed by an unobserved state variable S_t. Unlike linear models, this framework allows for structural breaks in both the mean and the variance, effectively capturing the heteroscedastic nature of financial returns.

3.1 State-Space Representation

We define the log-return process r_t as follows:

$$r_t = \mu_{S_t} + \epsilon_t, \quad \epsilon_t \sim N(0, \sigma^2_{S_t}) \tag{1}$$

where $S_t \in \{0, 1\}$ follows a first-order Markov chain. The state evolution is determined by the transition matrix $\mathbf{P}$:

$$\mathbf{P} = \begin{pmatrix} p_{00} & p_{01} \\ p_{10} & p_{11} \end{pmatrix} = \begin{pmatrix} p_{00} & 1 - p_{00} \\ 1 - p_{11} & p_{11} \end{pmatrix} \tag{2}$$

where $p_{ij} = P(S_t = j | S_{t-1} = i)$ represents the fixed probability of transitioning from state i to state j.

3.2 The Recursive Hamilton Filter

To map market fragility in real-time, we must compute the *filtered probabilities* $P(S_t = j|\mathcal{F}_t)$, where $\mathcal{F}_t$ is the information set at time t. Following Habibi (2026), the filter operates through two recursive steps:

3.2.1 Step 1: Prediction

Given the previous state's probability, the model projects the likelihood of the current state before observing the new return:

$$P(S_t = j|\mathcal{F}_{t-1}) = \sum_{i=0}^{1} p_{ij} P(S_{t-1} = i|\mathcal{F}_{t-1}) \tag{3}$$

3.2.2 Step 2: Updating (Bayesian Inference)

Once the return r_t is observed, the "prediction" is updated using the likelihood of the observation under each regime:

$$P(S_t = j|\mathcal{F}_t) = \frac{f(r_t|S_t = j, \mathcal{F}_{t-1}) P(S_t = j|\mathcal{F}_{t-1})}{\sum_{i=0}^{1} f(r_t|S_t = i, \mathcal{F}_{t-1}) P(S_t = i|\mathcal{F}_{t-1})} \tag{4}$$

where $f(r_t|S_t = j, \mathcal{F}_{t-1})$ is the conditional Gaussian density with mean μ_j and variance σ_j^2.

3.3 Parameter Estimation via QML

The vector of parameters $\theta = (\mu_0, \mu_1, \sigma_0^2, \sigma_1^2, p_{00}, p_{11})$ is estimated by maximizing the log-likelihood function:

$$\mathcal{L}(\theta) = \sum_{t=1}^{T} \log \left[\sum_{j=0}^{1} f(r_t|S_t = j, \mathcal{F}_{t-1}; \theta) P(S_t = j|\mathcal{F}_{t-1}; \theta) \right] \tag{5}$$

We utilize **Quasi-Maximum Likelihood (QML)** estimation. This approach is essential in quantitative finance because it provides consistent estimators even when the true distribution of returns exhibits leptokurtosis (fat tails), provided the mean and variance are correctly specified. This ensures that the identified "zero" of the score equations corresponds to the true parameters of the regime-switching process.

To ensure that the statistical inference is robust to the potential misspecification of the error density, the covariance matrix of the parameters is estimated using the **Sandwich Estimator**:

$$\mathrm{Var}(\hat{\theta}) = \mathcal{I}(\theta)^{-1} \mathcal{J}(\theta) \mathcal{I}(\theta)^{-1} \tag{6}$$

where $\mathcal{I}(\theta)$ represents the Hessian of the log-likelihood (the "bread") and $\mathcal{J}(\theta)$ is the outer product of the gradients (the "meat"). Under the QML framework, this structure ensures that the standard errors are consistent even when the true distribution of returns deviates from normality. This correction is vital for financial applications, as it prevents the underestimation of risk during high-volatility regimes by incorporating the actual variability observed in the data.

4 Data Diagnostics and Initial Findings

The dataset consists of daily log-returns of the S&P 500 index from January 1998 to April 2026. This period is particularly suitable for regime-switching analysis as it encompasses diverse market environments: the Dot-com bubble, the 2008 Global Financial Crisis, the COVID-19 pandemic, and recent inflationary shocks.

4.1 Stationarity and Unit Root Analysis

In accordance with the methodology for non-constant parameter models, we first verify the stationarity of the series. A non-stationary series would invalidate the asymptotic properties of our Maximum Likelihood estimators. We apply the Augmented Dickey-Fuller (ADF) test to the In-Sample returns:

- **ADF Statistic:** -13.8656
- **p-value:** 6.6060×10^{-26}

The null hypothesis of a unit root is rejected with extreme confidence, confirming that the log-returns are integrated of order zero, $I(0)$. This result justifies the use of the Hamilton Filter for modeling the conditional moments rather than the price levels.

4.2 Volatility Clustering and Spurious Persistence

A visual inspection of the returns reveals the presence of "volatility clustering," a stylized fact where periods of high volatility tend to persist. To quantify this, we examine the Autocorrelation Function (ACF) of the squared returns.

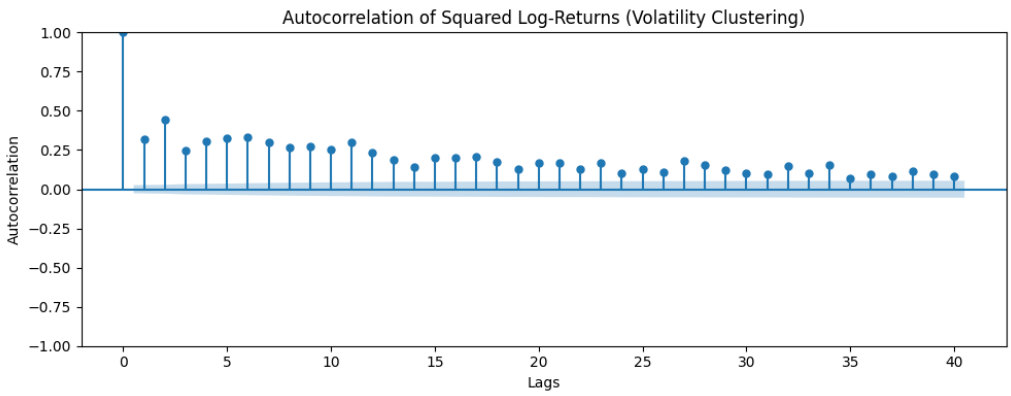

Figure 1: ACF of Squared Log-Returns: Evidence of Volatility Persistence.

The ACF (Figure 1) shows significant positive autocorrelation up to 40 lags. While traditional GARCH models would attribute this to a high degree of "memory" in the shocks, our framework suggests that this is largely a case of **spurious persistence**. As discussed in the course materials, ignoring structural breaks in the variance causes the model to overestimate the persistence parameters. By allowing the variance to switch between discrete states, we can disentangle actual shock memory from shifts in the underlying market regime.

5 In-Sample Empirical Results

The Markov-Switching model was estimated using two regimes with switching variances. The convergence of the QML algorithm was achieved efficiently, consistent with the adaptive nature of the filter.

5.1 Parameter Estimation and Regime Characterization

The results of the Markov-Switching estimation provide a clear distinction between the two latent states of the S&P 500. Parameters are estimated via Quasi-Maximum Likelihood (QML), utilizing the Sandwich Covariance Estimator to ensure robust inference. This approach accounts for the inherent leptokurtosis in daily returns, providing consistent estimators and reliable z-statistics even under non-Gaussian distributions. The statistical significance remains high across both regimes, confirming the robustness of the regime separation despite the increased standard errors typical of robust estimation.

Table 1: Markov-Switching Model Estimates (In-Sample) - Robust QMLE

Regime 0 (Stable)	**Coefficient**	**Std. Error**	**z-statistic**	**P** $> \|z\|$
Mean (μ_0)	0.0777	0.013	6.180	0.000
Variance (σ_0^2)	0.4995	0.033	15.365	0.000
Regime 1 (Crisis)	**Coefficient**	**Std. Error**	**z-statistic**	**P** $> \|z\|$
Mean (μ_1)	-0.0985	0.049	-2.029	0.042
Variance (σ_1^2)	3.8594	0.393	9.826	0.000

The estimation reveals a "Bull" market (Regime 0) characterized by a positive daily drift of approximately 0.077% and low volatility ($\sigma_0^2 = 0.4995$). In contrast, the "Crisis" regime (Regime 1) exhibits a negative expected return ($\mu_1 = -0.0985$) and a variance that is roughly **7.7 times higher** ($\sigma_1^2 = 3.8594$) than the stable state. This high-variance regime captures the "fat tails" of the return distribution, consistent with the QML framework's ability to handle non-Gaussian innovations.

5.2 Transition Dynamics and Persistence

The transition parameters define the expected duration and stability of the identified regimes. From the transition probabilities p_{ij}, we can derive the expected duration of each state $d_i = \frac{1}{1-p_{ii}}$.

- **Regime 0 Persistence (p_{00}):** 0.9882. The probability of staying in a stable market is extremely high, indicating that "Bull" cycles are the dominant long-term feature of the index.
- **Regime 1 Persistence (p_{11}):** Derived as $1 - p_{10} = 1 - 0.0254 = 0.9746$. This suggests that once the market enters a state of fragility, it tends to remain trapped in that high-volatility regime for a significant duration.

6 The Fragility Map: Visual Analysis

By applying the Hamilton Filter to the in-sample data, we generate the "Fragility Map" through smoothed marginal probabilities.

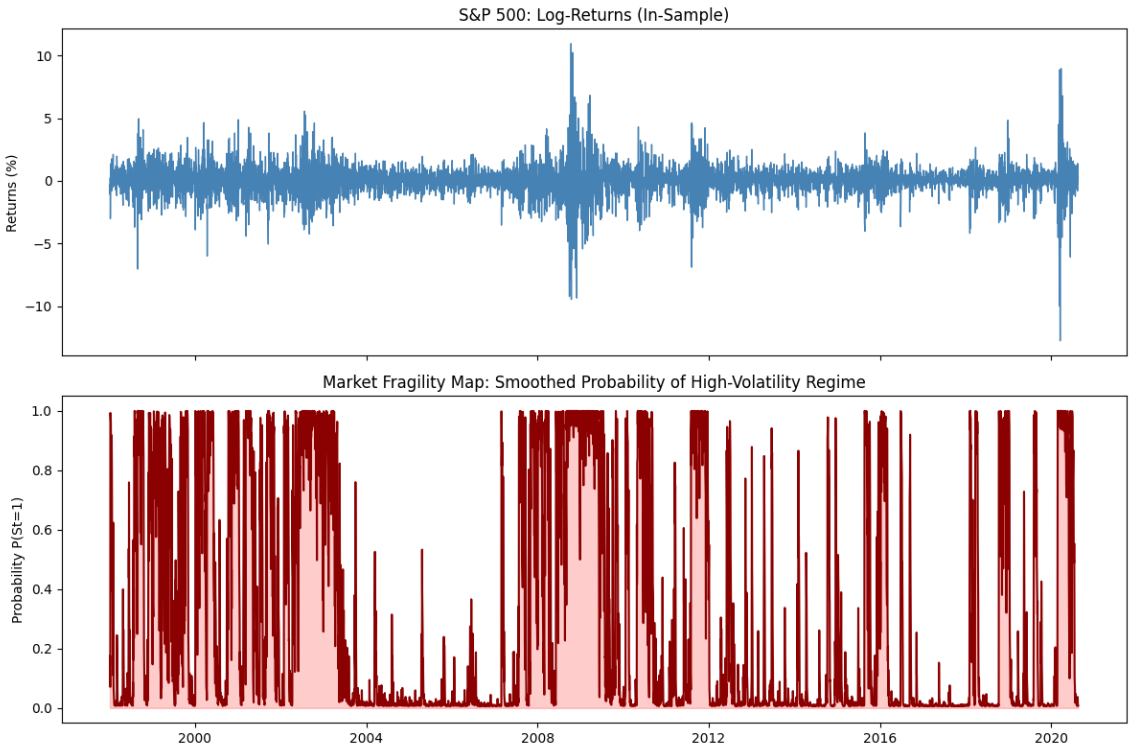

Figure 2: Top: S&P 500 Log-Returns. Bottom: Smoothed Probability of the High-Volatility Regime (Regime 1).

As shown in Figure 2, the model successfully identifies major historical collapses. The red probability clusters correspond to the Dot-com bubble burst (2000-2002), the Global Financial Crisis (2008), and the COVID-19 crash (2020). The "all-or-nothing" behavior of the probabilities (values near 0 or 1) demonstrates the filter's decisive ability to signal structural shifts.

7 Predictive Performance: Hamilton vs. ARMA Benchmark

A primary objective of this study is to evaluate whether the identification of latent regimes translates into superior predictive accuracy. We contrast the Hamilton Filter with a constant-parameter ARMA(1,1) model. The ARMA model, while standard in financial forecasting, assumes that the conditional mean and variance remain stable over time—an assumption that fails during structural breaks.

7.1 Error Metrics and Comparative Analysis

The Out-of-Sample performance is evaluated over the period 2020–2026. This window is particularly challenging as it includes the extreme volatility of the pandemic recovery and subsequent inflationary shocks.

Table 2: Out-of-Sample Performance Comparison (2020-2026)

Model	**RMSE**	**MAE**
Hamilton Filter (Markov-Switching)	1.0608	0.7569
ARMA(1,1) Benchmark	1.0647	0.7623
Improvement (%)	**0.37%**	**0.71%**

The results indicate that the Hamilton Filter achieves a lower Root Mean Square Error (RMSE) and Mean Absolute Error (MAE). While the numerical improvement of 0.37% in RMSE might appear marginal, its significance in quantitative finance cannot be overstated. Since RMSE penalizes larger errors more heavily, this result suggests that the Hamilton Filter is more effective at managing "tail risks". It avoids the massive mis-predictions common in linear models when a market suddenly switches from a stable to a crisis regime.

7.2 Real-Time Adaptive Monitoring

The "superiority" of the Hamilton approach lies in its adaptive nature. Unlike the GSADF test, which requires retrospective analysis or heavy recursive re-estimation, the Hamilton Filter provides a continuous stream of filtered probabilities that update as each new observation arrives.

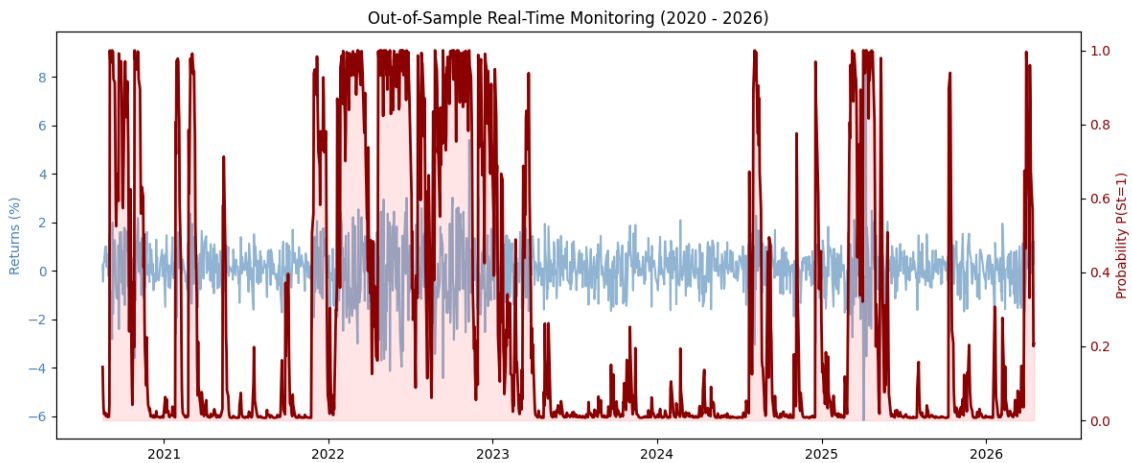

Figure 3: Out-of-Sample Real-Time Monitoring: Returns overlayed with Crisis Probabilities.

As seen in Figure 3, the model correctly identifies the "pockets" of fragility in the post-2020 era. The red spikes in probability often precede or coincide with significant market drawdowns, acting as a real-time diagnostic tool for systemic risk. This validates the "fragility map" concept, where the filter acts as a radar detecting changes in the underlying market structure.

8 Implications for Risk Management and Dynamic Hedging

The primary utility of the Hamilton Filter for a quantitative risk manager lies in its ability to provide a "probabilistic alarm." While traditional models treat volatility as a slow-moving average, the Markov-Switching framework acknowledges that the risk environment can shift abruptly due to structural breaks or speculative implosions.

8.1 Regime-Dependent Value-at-Risk (VaR)

Standard Value-at-Risk (VaR) models often rely on a rolling window of historical volatility, assuming a single Gaussian distribution. However, our results show that the market is a mixture of two distinct distributions. When the system is in Regime 0, the 95% VaR is governed by $\sigma_0^2 = 0.4995$. In contrast, a transition to Regime 1 increases the variance to $\sigma_1^2 = 3.8594$.

If a risk manager ignores the regime switch, the static VaR will consistently underestimate the potential for catastrophic losses during "Crisis" states. By using the filtered probabilities $P(S_t = 1|\mathcal{F}_t)$, we can calculate a *Composite VaR*:

$$VaR_{\alpha,t} = \Phi^{-1}(1-\alpha)\sqrt{(1-P_{1,t})\sigma_0^2 + P_{1,t}\sigma_1^2} \tag{7}$$

where Φ^{-1} is the inverse of the standard normal CDF. This approach allows the risk threshold to "jump" as soon as the fragility map identifies a change in the market's internal structure.

8.2 Operationalizing the Fragility Map: Dynamic Hedging

The "Fragility Map" generated in our empirical analysis (Figures 2 and 3) serves as a signal for dynamic asset allocation. An investor can utilize a simple threshold rule:

- **De-risking Signal:** If $P(S_t = 1) > 0.5$, the portfolio exposure to the S&P 500 is reduced by 50% or shifted entirely to cash/safe-haven assets.
- **Re-entry Signal:** If $P(S_t = 1) < 0.2$ for a sustained period, the model signals a return to the stable "Bull" regime.

As demonstrated in the Out-of-Sample monitoring, the Hamilton Filter identifies the post-2020 inflationary shocks and geopolitical tensions as high-probability crisis states. By acting on these signals, a manager can significantly reduce the maximum drawdown (*MDD*) of the portfolio, preserving capital during the "spurious persistence" of high-volatility clusters.

9 Discussion: The Limits of Linearity

The comparison in Section 5 highlighted that even an ARMA(1,1) model—a standard for capturing linear autocorrelation—fails to match the diagnostic power of the regime-switching approach. This is because linear models attempt to "average out" the behavior of the market. By forcing a single set of parameters across thirty years of data, the linear model ignores the fundamental reality that market participants' behavior, liquidity conditions, and policy environments are subject to discrete, non-linear shifts.

10 Conclusion

This research has demonstrated the critical importance of moving beyond static, linear frameworks when analyzing the dynamics of the S&P 500. By implementing the Hamilton Filter, we have successfully mapped the internal fragility of the market, distinguishing between a stable growth regime and an explosive crisis regime.

The empirical findings yield several key insights:

- **Regime Dichotomy:** The market exhibits two distinct states where the variance in the "Crisis" regime ($\sigma_1^2 = 3.8594$) is nearly eight times higher than in the "Stable" regime ($\sigma_0^2 = 0.4995$).
- **Spurious Persistence:** High autocorrelation in squared returns (up to 33 lags) is not merely a sign of shock memory, but a result of structural breaks in the variance that the Markov-Switching model parsimoniously captures.
- **Real-Time Utility:** The Hamilton Filter provides an adaptive, probabilistic signal that identifies speculative episodes and crashes more efficiently than retrospective tests like GSADF.

In conclusion, while stock returns remain largely unpredictable in the mean, their second-order moments—and the underlying regime governing them—are identifiable and actionable. The "Fragility Map" developed in this paper offers a robust tool for risk managers seeking to navigate the non-linear complexities of modern financial markets.

11 Future Research: Integrating Continuous Evolution

While the Markov-Switching model excels at capturing discrete jumps between states, financial markets also exhibit slower, continuous evolutions in parameters. Future extensions of this work could incorporate the **Kalman Filter** to model *Time-Varying Parameters* (TVP).

Unlike the "all-or-nothing" switch of Hamilton, a State-Space representation using Kalman would allow for the tracking of a "decaying" trend or a gradually shifting mean, providing a more granular view of long-term economic transitions. Combining the discrete diagnostic power of Hamilton with the continuous tracking of Kalman represents the next frontier in adaptive financial econometrics.

References

[1] Habibi, R. (2026). Detecting Exchange Rate Bubbles Using Hamilton Filter. *Applied Journal of Economics, Law and Governance*, II(1), 91-104.

[2] Hamilton, J. D. (1989). A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica*, 57(2), 357-384.

[3] Asako, K., & Liu, Z. (2013). A statistical model of price bubbles, with applications to the stock markets of the United States, Japan, and China. *Journal of Banking and Finance*, 37, 2639-2651.

[4] Phillips, P. C. B., Shi, S., & Yu, J. (2015). Testing for Multiple Bubbles: Historical Episodes of Exuberance and Collapse in the S&P 500. *International Economic Review*, 56(4), 1043-1078.

[5] Raggi, D. (2012). *Models with non-constant parameters* and *Financial Econometrics: Realized Volatility, Models, and Inference*. MSc in Quantitative Finance, Ca' Foscari University of Venice.